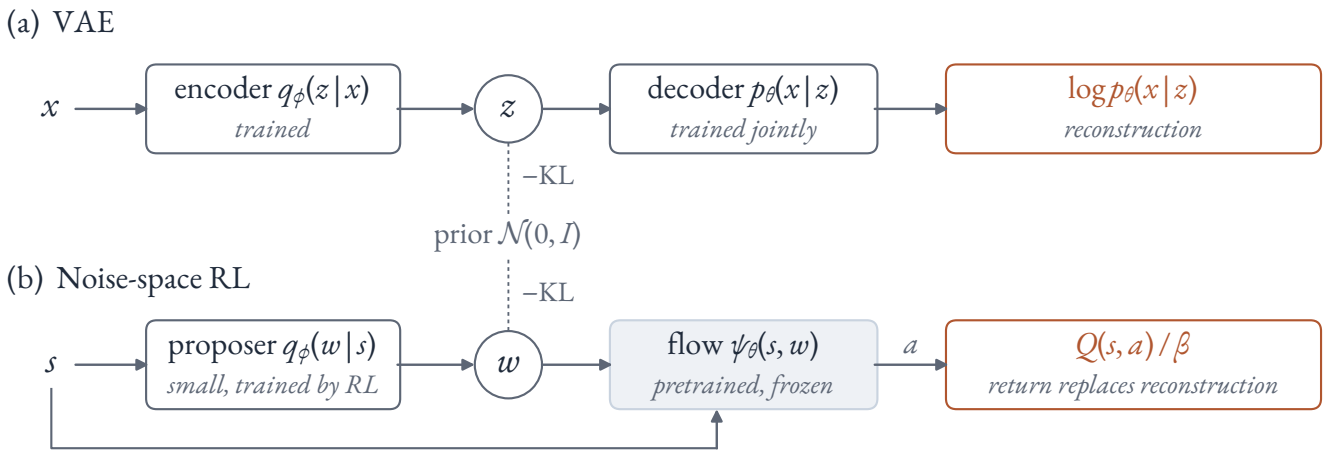
(a) VAE
x
encoder $q_\phi(z \mid x)$
trained
z
decoder $p_\theta(x \mid z)$
trained jointly
$\log p_\theta(x \mid z)$
reconstruction
−KL
prior $\mathcal{N}(0, I)$
−KL
(b) Noise-space RL
s
proposer $q_\phi(w \mid s)$
small, trained by RL
w
flow $\psi_\theta(s, w)$
pretrained, frozen
a
$Q(s, a) / \beta$
return replaces reconstruction