

YUHAN CHI

School of Mathematical Sciences, Fudan University

yhchi25@m.fudan.edu.cn · chi-shan0707.github.io · github.com/Chi-Shan0707

RESEARCH INTERESTS

Reinforcement learning and generative models; optimization and control for decision-making under uncertainty.

EDUCATION

Fudan University

B.S. in Information and Computing Science & Artificial Intelligence (Double Degree)

GPA 3.98/4.00 · Rank 3/228

Shanghai, China

Sep. 2025 – Present

PUBLICATIONS

Calibrating What Behavior Can Identify: Conformal Reward Sets for Inverse Reinforcement Learning

Yuhan Chi. *NeurIPS 2026 Workshop on Physical Understanding for Decision-Making (PUDM)*, Poster [Paper](#) · [Code](#) · [Project](#)

- Adapts conformal inverse optimization to IRL with a coordinate-invariant, policy-value-based score for calibrated reward sets.
- Characterizes when behavioral compatibility implies true-reward coverage, and proves when robust planning over the set cannot certify improvement over a reference policy.

Token-Level Verification under Controlled Evaluation

Yuhan Chi. *ICML 2026 Workshop on AI for Math (AI4Math)*, Poster

[Paper](#) · [Code](#)

- Shows that evaluation-protocol choices alone shift apparent verification AUROC of LLM math traces by up to **0.18**.
- Under controlled scoring, shallow token statistics reach only **0.60–0.75** AUROC; final-token entropy drops from 0.72–0.75 to chance (0.47–0.48).

RESEARCH EXPERIENCE

Visiting Student, Mechanical Systems Control Lab

University of California, Berkeley

Berkeley, CA

Jul. 2026 – Aug. 2026

- Studied robot hardware interfaces and explored dexterous-hand grasping in simulation.

SELECTED PROJECTS

code-not-text — cross-domain measurement study, 31,040 reasoning runs

[Demo](#)

- Selecting among 64 candidates with cheap reasoning-trace features changes pass@1 by **+10.0 pp** (math), **+8.0 pp** (science), and **−0.6 pp** (code); rules out five alternative explanations for the coding failure.

TinyLoRA-GRPO-Coder — RL post-training for code generation

- Trains 32 shared scalars of Qwen2.5-Coder-3B with GRPO, rewarded by real g++ compile-and-run outcomes.

HONORS, TALKS & SERVICE

Reviewer, ICML 2026 Workshop on AI for Math (AI4Math)

2026

Talk, *Reinforcement Learning: From Bandits to PPO*, Fudan University

Apr. 2026

First Prize, National Olympiad in Informatics in Provinces (NOIP), Shanghai

2023

SKILLS

Programming Python, C++, C

ML PyTorch, Machine Learning, Reinforcement Learning, Generative Models

Tools Git, Linux, LaTeX

Languages Chinese (native), English (fluent)

SELECTED COURSEWORK

Mathematical Analysis I, II (A, A+) · Advanced Algebra I, II (A−, A) · Analytic Geometry (A) · Programming (A) · Introduction to AI (A) · Foundations of Software for AI (A)

OPEN SOURCE & COMMUNITY

Founder, Fudan Open Source Initiative

Projects: FDUGuideBook, FDU-Sharing, ic-guide, github-unflag-playbook-cn